JONNY HOFMEISTER

Tanzanian Waterpoint Classification

Predicting Waterpoint Functionality

Presentation Overview

01

Business Probelm

02

Data

03

Modeling Results
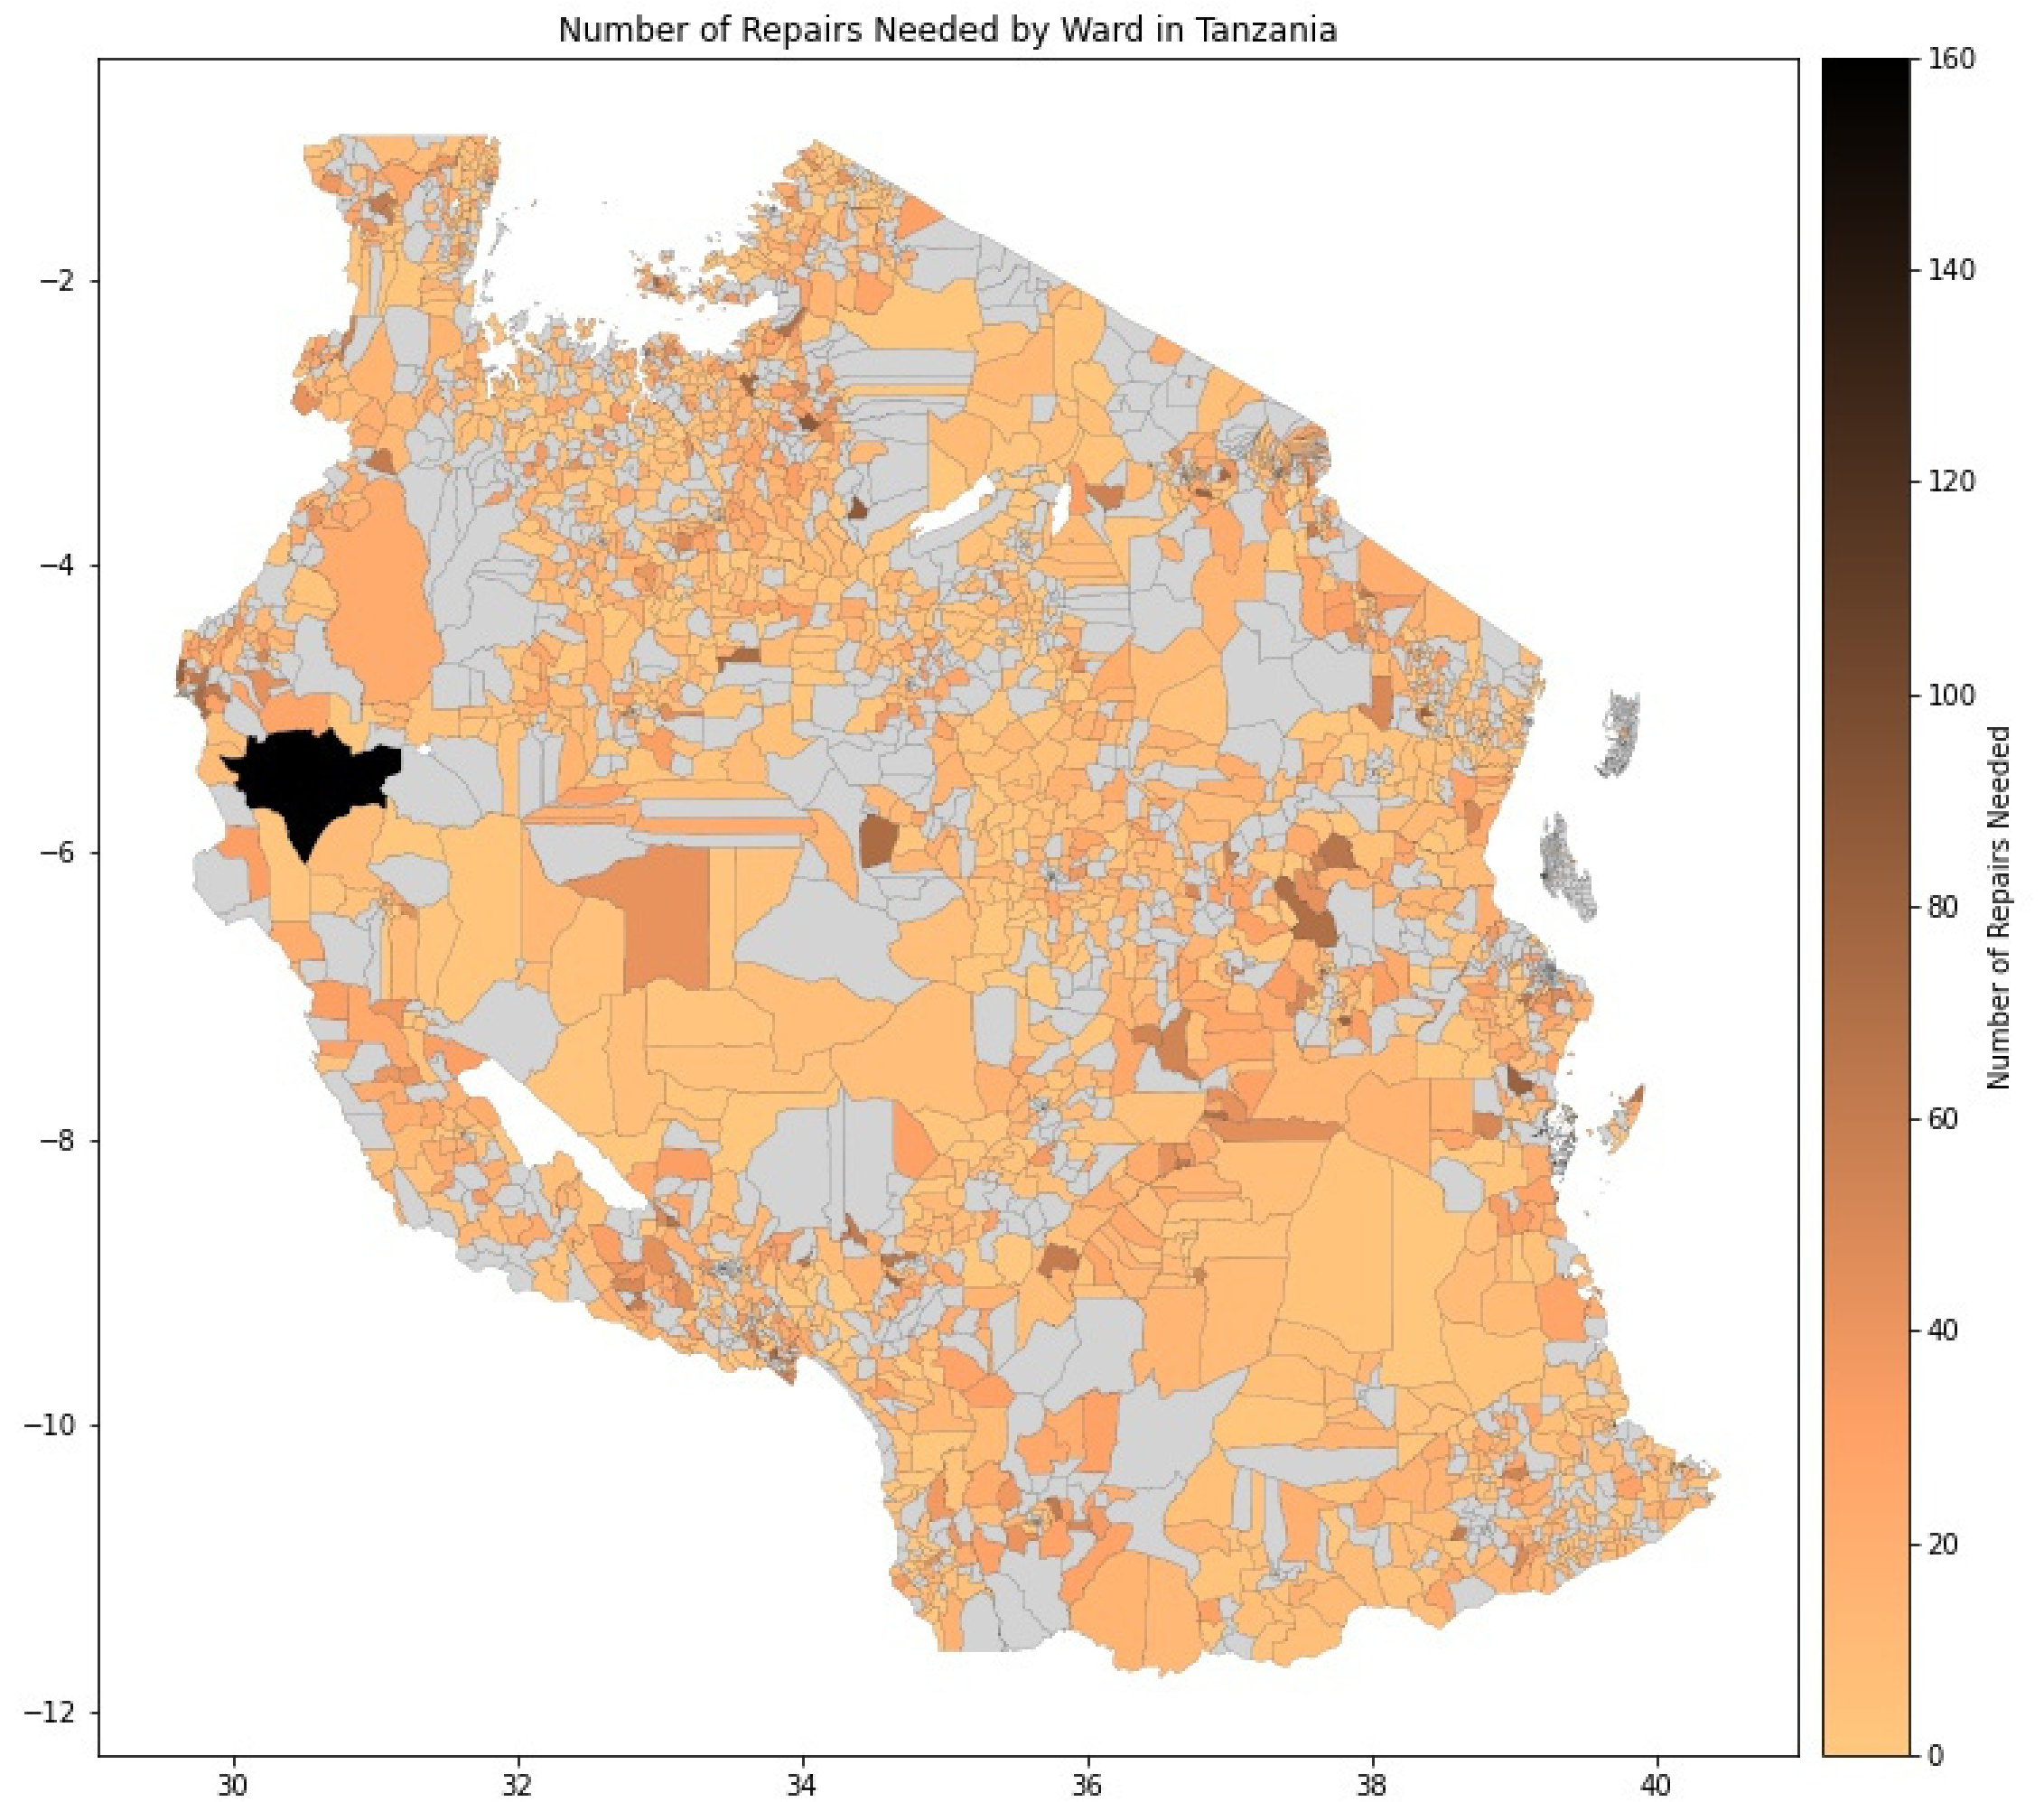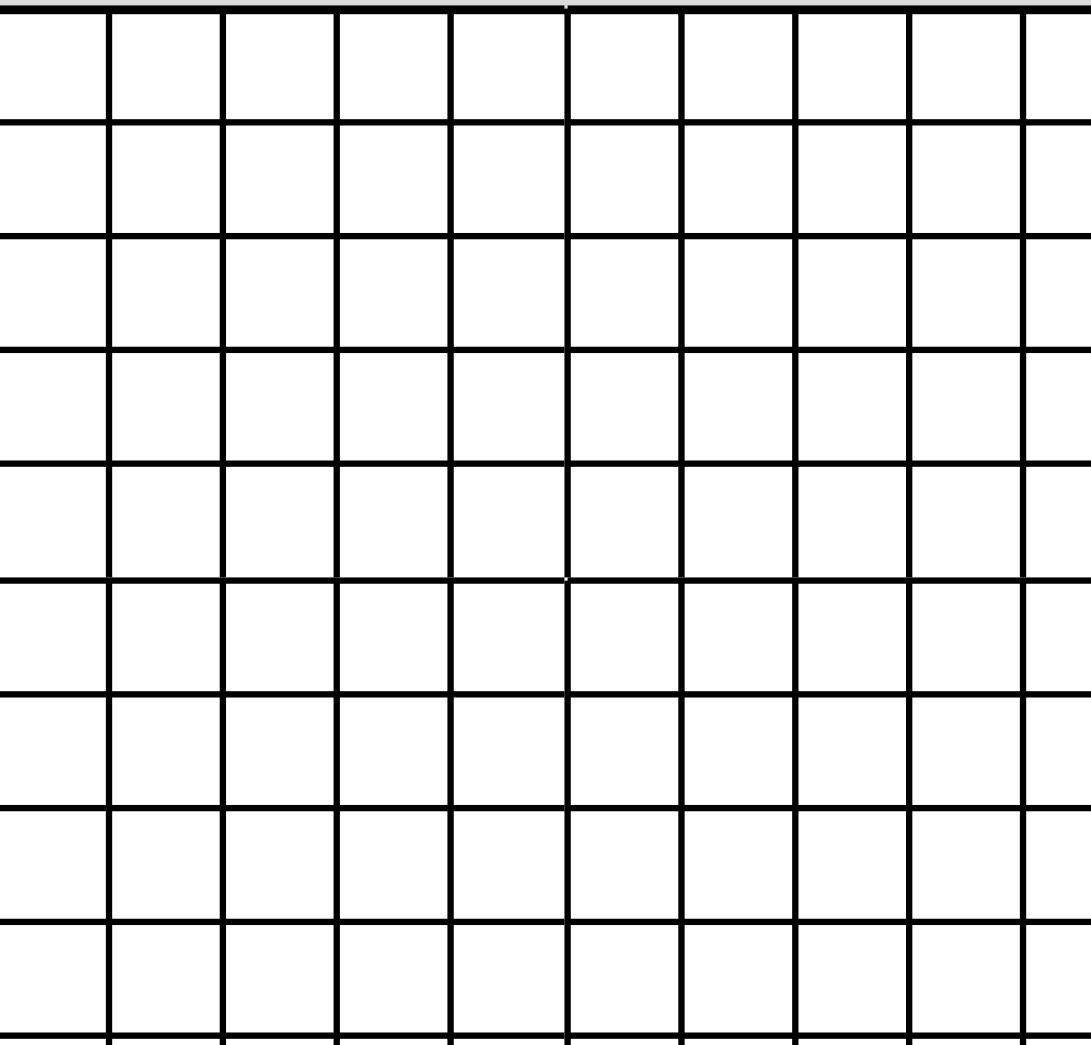
04

Moving Forward

Business Problem

Predict Functional Status of Waterpoints

- Train a model on waterpoints with known functionality.
- Predict the functionality of unknown waterpoints across Tanzania.
- Then, determine best practices for addressing repair.

Map

Shows the extent of missing data

Number of Repairs Needed by Ward in Tanzania

Data Understanding

Target: Functional Status Group

Ternary Classification

- 'Functional'
- 'Non Functional'
- 'Functional Needs Repair'

Metrics:

How to analyze classification success

In the context of efficiently performing repairs, a FN isn't any worse than a FP. So balance Precision and Recall

Goal: Maximize F1 score

Confusion Matrices and Accuracy will also be used in evaluation.

Data Understanding:
Features

Numericals: Total Static Head, Population, Year Created, etc

Booleans: Public Meeting, Permit

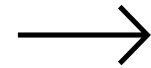**Categoricals:** Can be OHE or Freq Encoded

OHE: Scheme Management, Extraction Type, Quantity, Source

Frequency: Location based columns - Region. LGA, Ward, etc

Scalers and Encoders:

- MinMax Scaler
- Simple Imputer
- One-Hot Encoder
- Count Encoder

Pipeline Process

End to End modeling with Sklearm →

01
Separate features

Sort columns by datatype

02
Impute missing values

Impute NaNs and zeros in categorical and numerical columns. Fill accordingly

03
Scale, One-Hot and Count Encode

Scale Numericals

OHE and Count (freq) encode categorical columns

04
Fit Models and Score

Fit the different model types.

Score them with desired metrics.

Examine confusion matrix.

05
Iterate Modeling

Iterate through modeling process.

Search for best parameters.

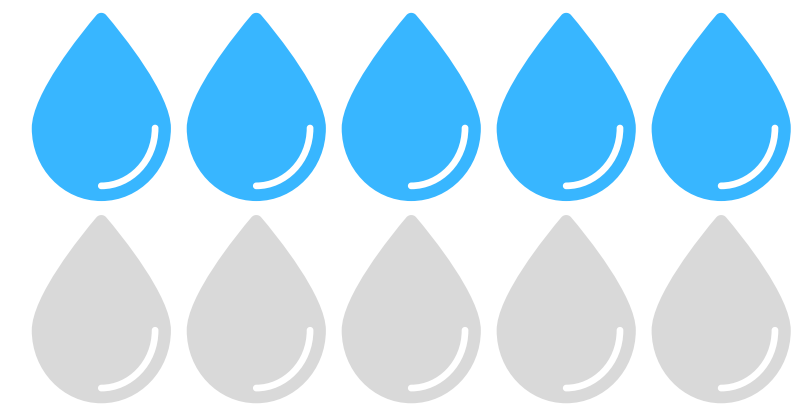Optimize the model for target metrics.

Repeat.

Baseline Model

Create a naive baseline to beat:

- If total static head > 0, predict 'functional'.
- If total static head == 0, predict 'non functional'

This at least predicts 2 classes, whereas predicting the majority class, 'functional', is not actionable
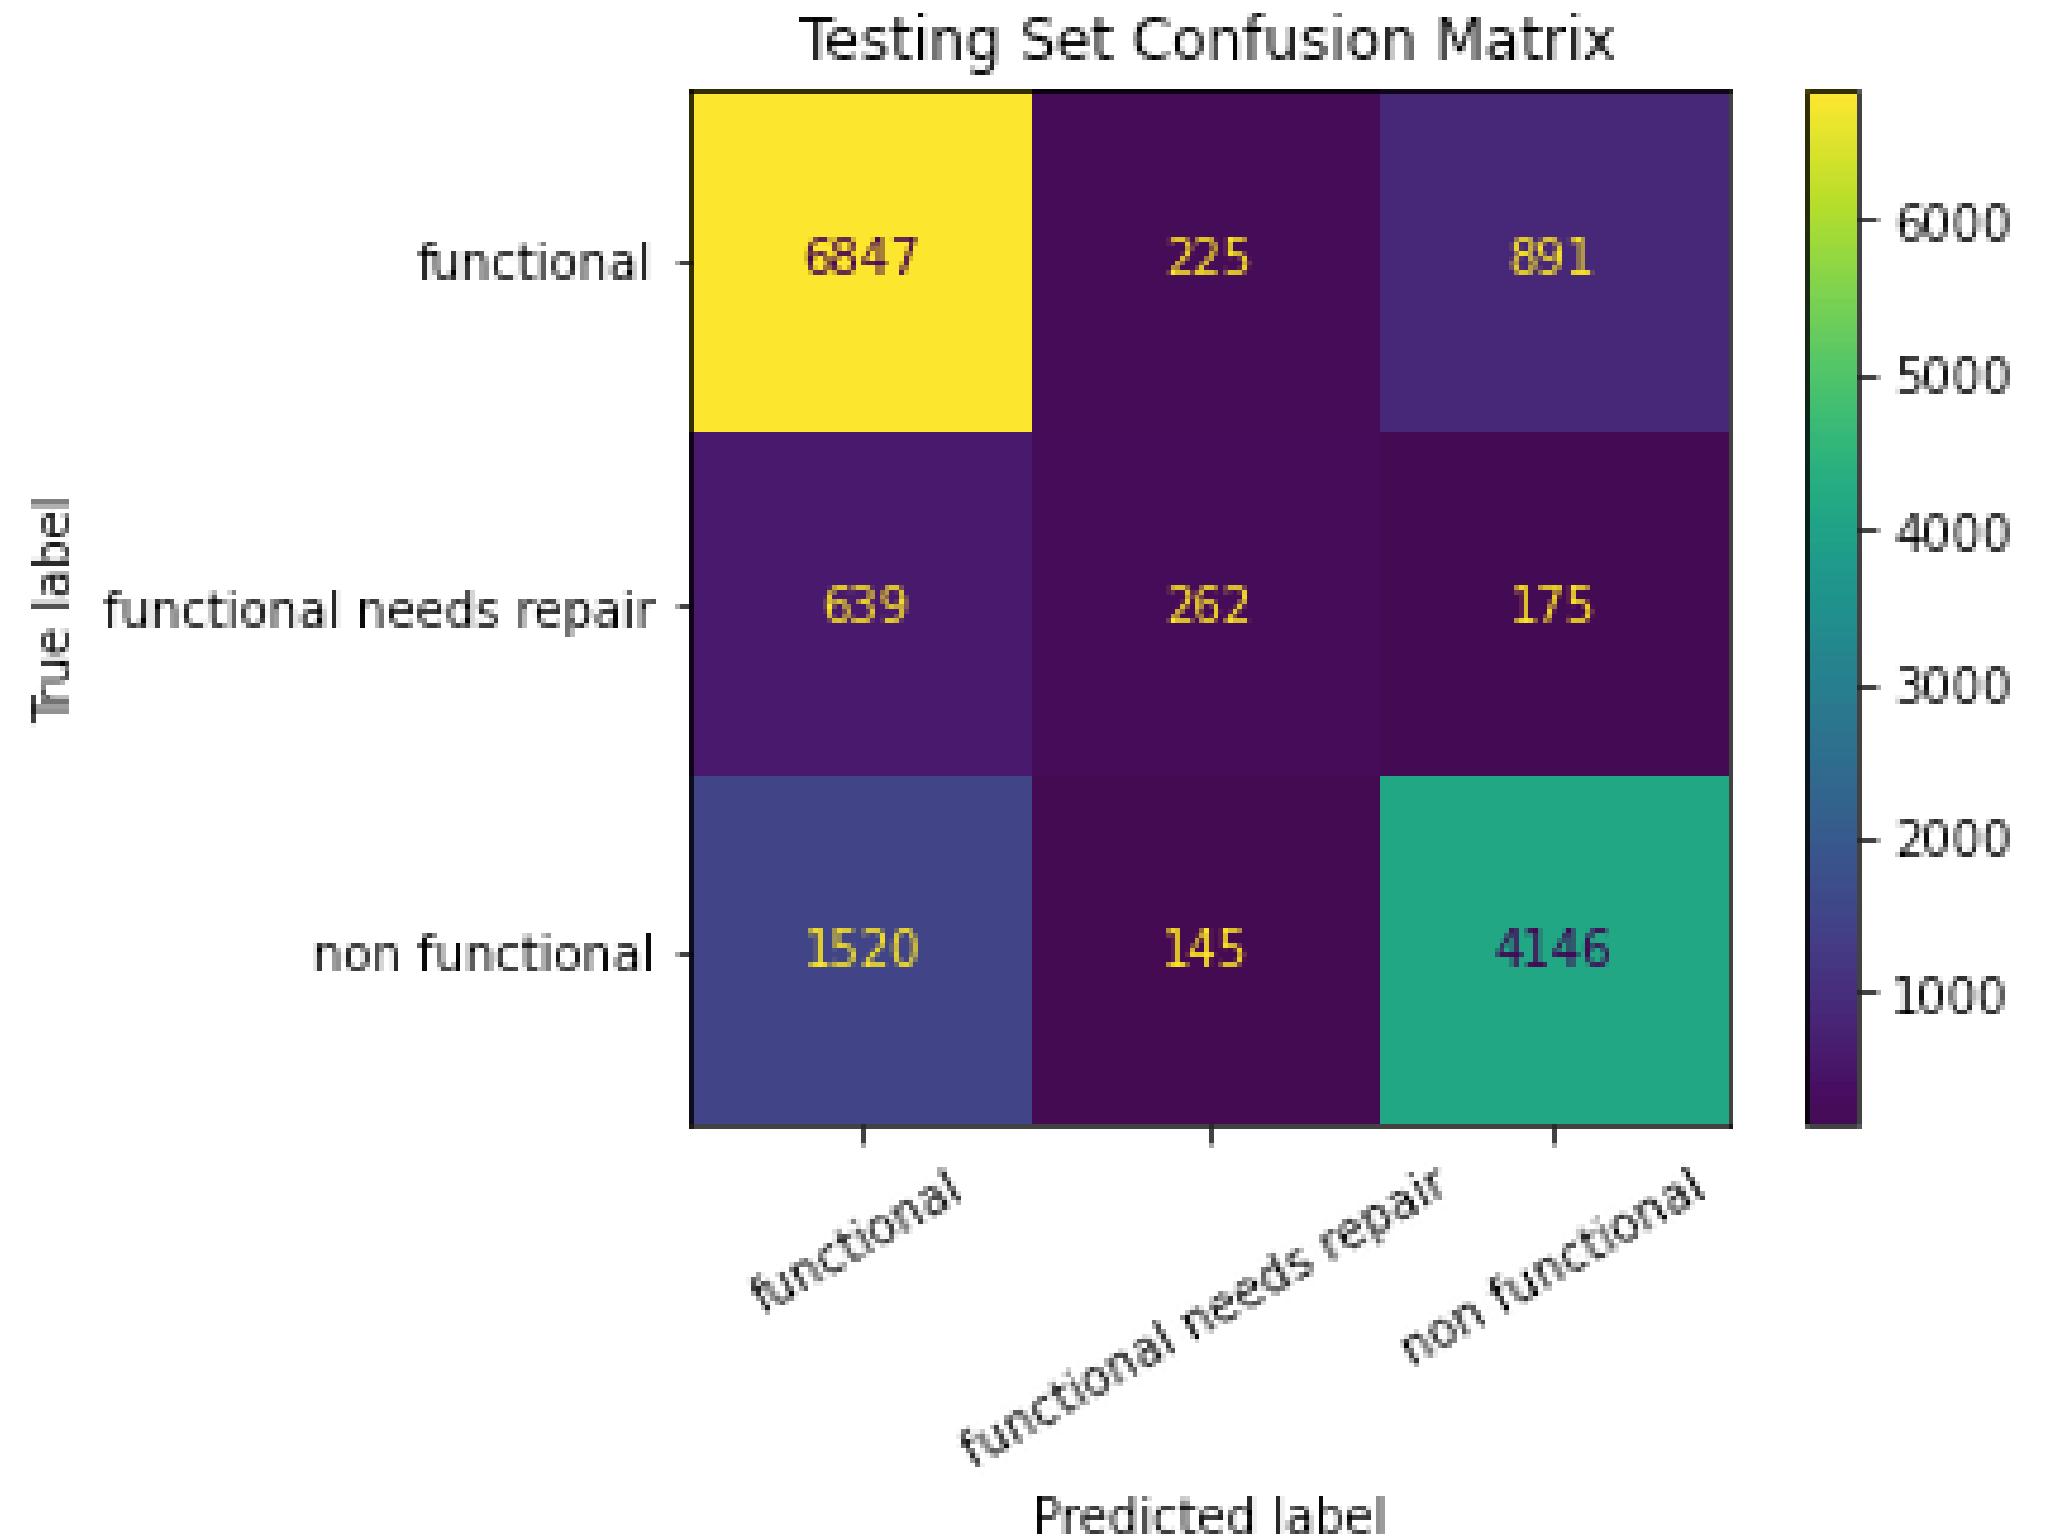
53% accurate

5 out of 10

Pumps predicted with correct functionality

K-Nearest Neighbors Model

Determines similarity between samples using distance metrics

- Training F1: 0.8
- Testing F1: 0.75
- Training Acc: 0.81
- Testing Acc: 0.76

Testing Set Confusion Matrix

Random Forest

SCORES BEFORE PARAMETER TUNING

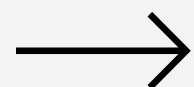Training F1: 0.96
Training Acc: 0.96

Testing F1: 0.76
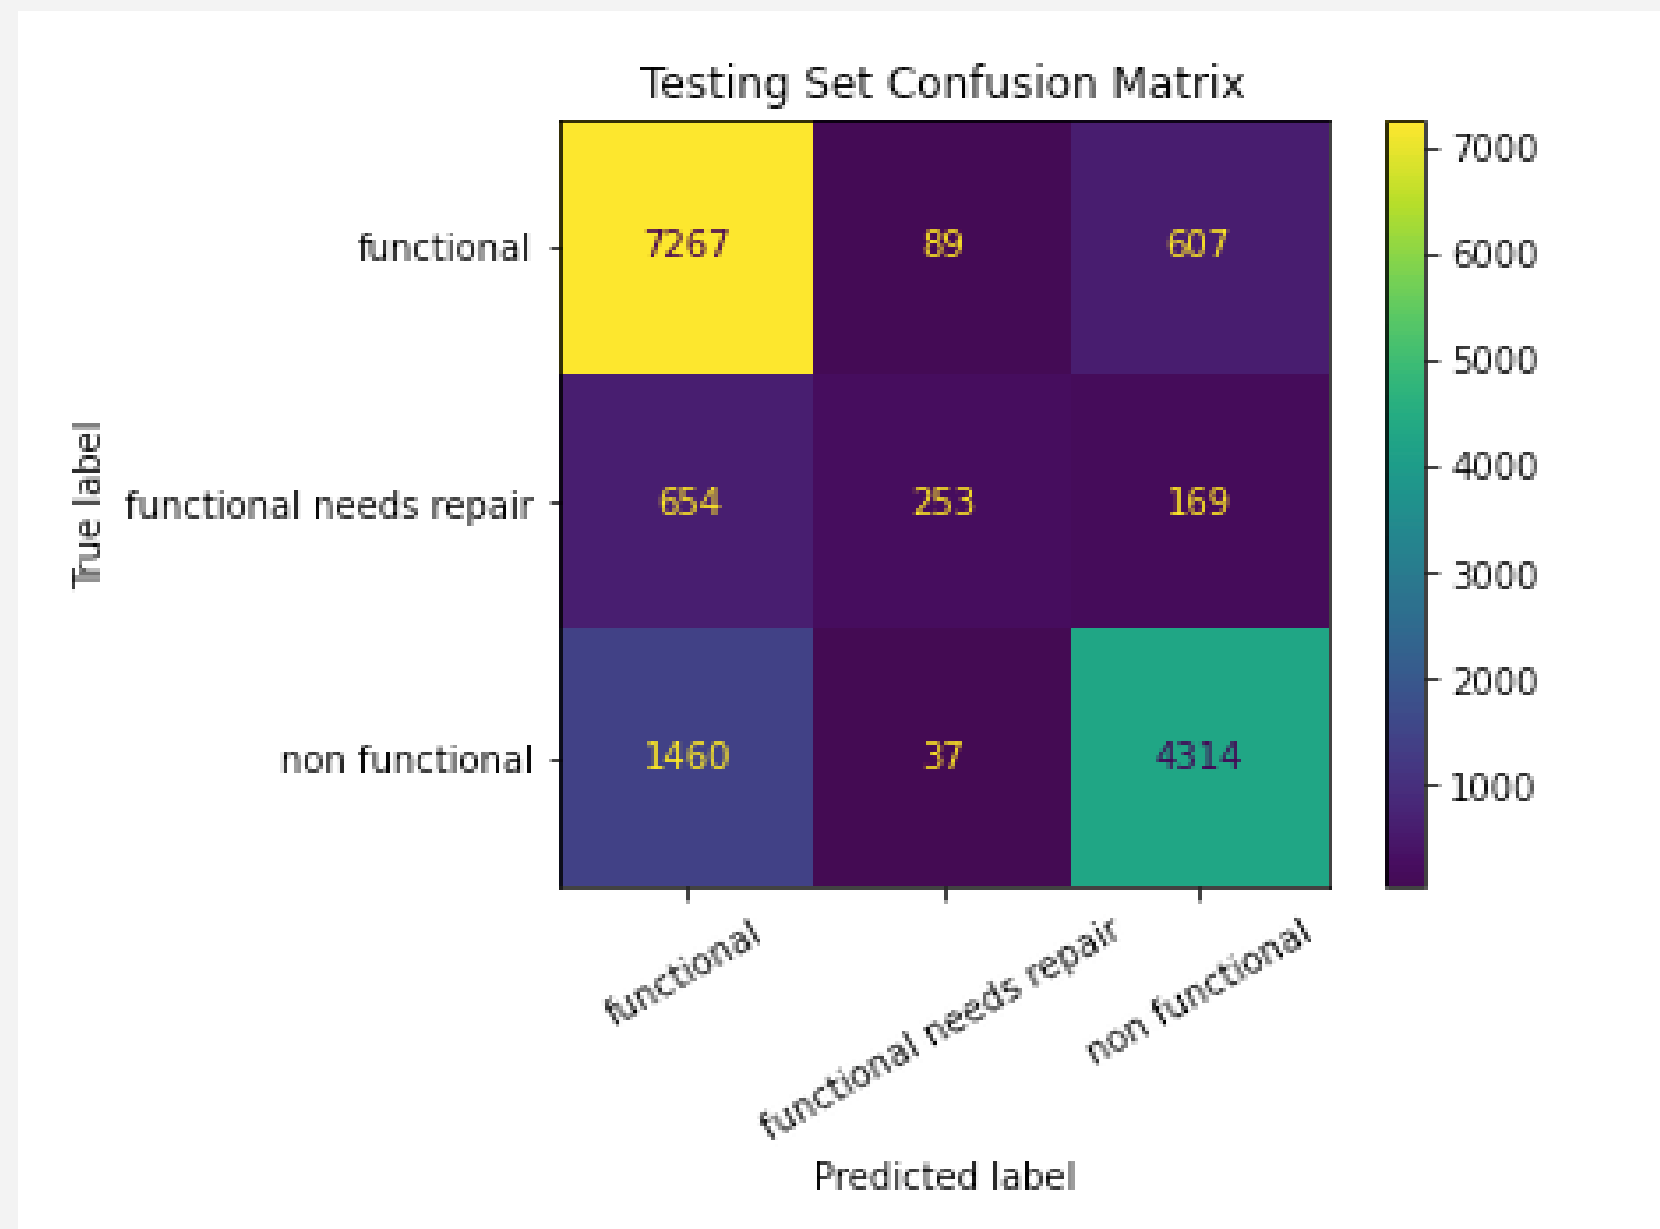Testing Acc: 0.78

GRID SEARCH

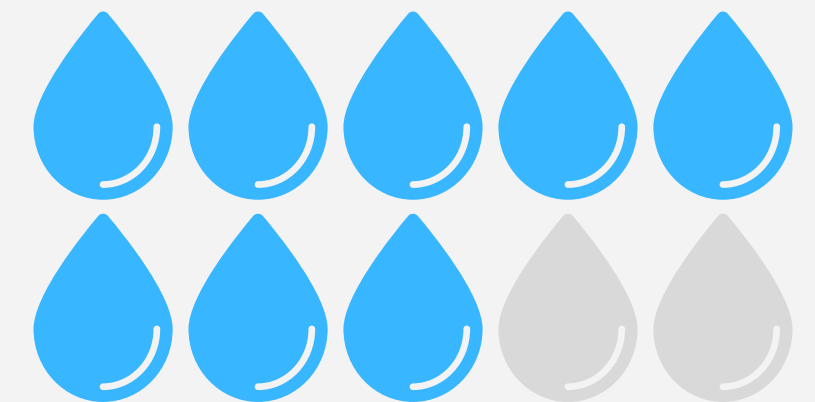Improved testing scores by 2% and reduced overfitting by 10%.

PARAMETERS

n_estimators = 160
max_depth = 30
min_samples_split = 2
min_samples leaf = 2

→

Decision trees created to classify samples with splits/if statements

Testing Set Confusion Matrix

	functional	functional needs repair	non functional
functional	7267	89	607
functional needs repair	654	253	169
non functional	1460	37	4314

True label / Predicted label

8 out of 10

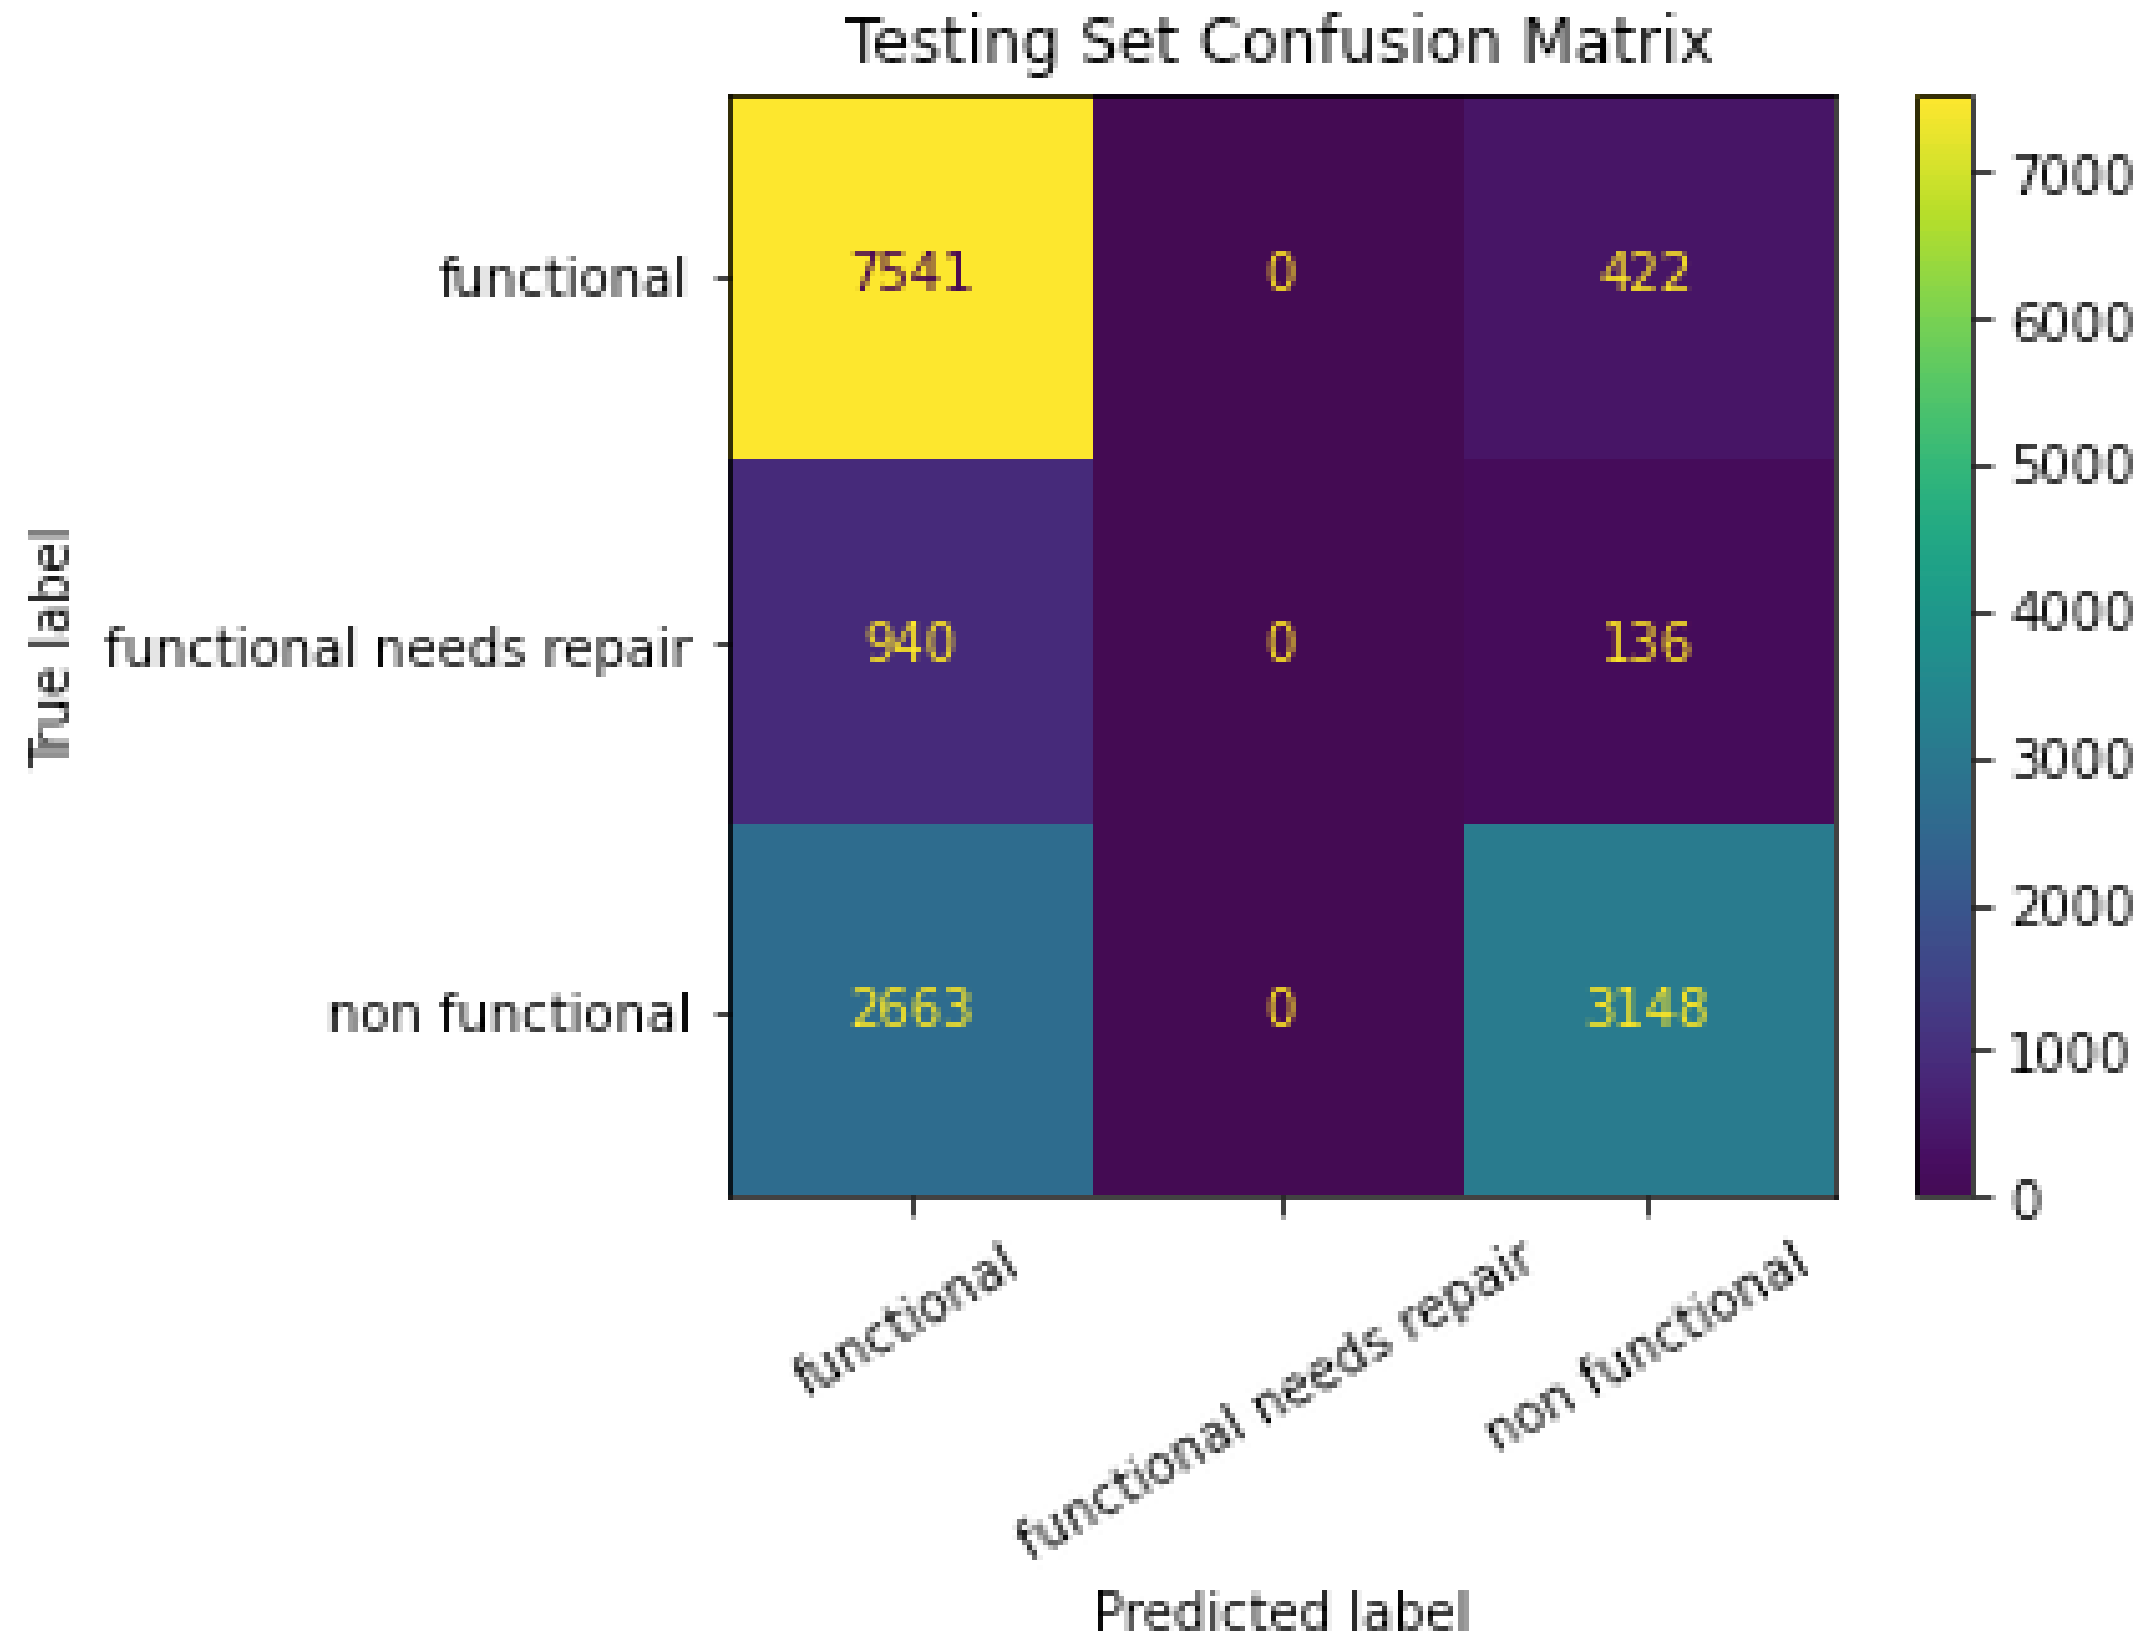Pumps predicted with correct functionality

0.78
Best testing F1

80%
Best testing accuracy

Support Vector Classifier

Manipulates parameter space to create optimized decision boundaries for classes

- Testing F1: 0.68

- Testing Acc: 0.72

- Failed to predict 3rd class

Testing Set Confusion Matrix

Results

Current Random Forest model can fill missing data with 80% accuracy and balanced Precision and Recall
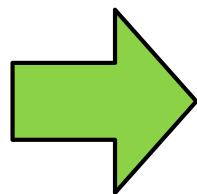
Next Steps

Rigorous Feature Selection

Search for features that produce signal and not noise.

If we want to predict functionality in wards with no data, perhaps location features wont add any information to the model, only overfit it.

Better Imputation Strategies

Rather than impute features like construction year with the median value, impute with a random selection from a probability distribution representing how often each year shows up in the data.
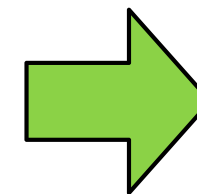
Repairs

Use the model to begin to fill in missing pump statuses in Tanzania.

Address the TSMs goals and constraints for repairs, begin to generate actionable insights for most efficient methods.

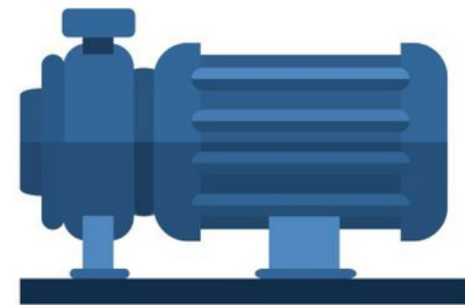Prioritize populated areas? Certain types of pumps? Etc.

Resource Page

Sources and Citations

Dataset on DrivenData competition:

DRIVENDATA

Pump fundamentals

Sklearn Documentation

Thank you for your time!

GitHub Link and email:
jonny.hofmeister@gmail.com